

# Projeto01.R

rodrigolima82

2019-05-19

```
# Script para checar as colunas do dataset
```

```
# Carregando os Pacotes
```

```
library(data.table)
```

```
# Carregando o Dataset #
```

```
# Os dados brutos contém 100.000 linhas e 8 colunas (atributos).
```

```
# A coluna "is_attributed" é o alvo.
```

```
dt <- fread("dados/train_sample.csv")
```

```
df <- as.data.frame(dt)
```

```
# Remove dt
```

```
rm('dt')
```

```
# Visualizando dados do dataframe
```

```
# View(df)
```

```
str(df)
```

```
## 'data.frame': 100000 obs. of 8 variables:
```

```
## $ ip : int 87540 105560 101424 94584 68413 93663 17059 121505 192967 143636 ...
```

```
## $ app : int 12 25 12 13 12 3 1 9 2 3 ...
```

```
## $ device : int 1 1 1 1 1 1 1 1 2 1 ...
```

```
## $ os : int 13 17 19 13 1 17 17 25 22 19 ...
```

```
## $ channel : int 497 259 212 477 178 115 135 442 364 135 ...
```

```
## $ click_time : chr "2017-11-07 09:30:38" "2017-11-07 13:40:27" "2017-11-07 18:05:24" "2017-11-07 18:05:24" ...
```

```
## $ attributed_time: chr "" "" "" "" ...
```

```
## $ is_attributed : int 0 0 0 0 0 0 0 0 0 0 ...
```

```
# Nome das variáveis
```

```
# ip, app, device, os, channel, click_time, attributed_time, is_attributed
```

```
# Aplicando Engenharia de Atributos em Variáveis Numericas
```

```
# Transformar o objeto de data
```

```
df$click_time <- as.POSIXct(df$click_time)
```

```
# Extraíndo dia e hora
```

```
df$click_day <- as.integer(format(df$click_time, "%d"))
```

```
df$click_hour <- as.integer(format(df$click_time, "%H"))
```

```
# Transformando variáveis numéricas em variáveis categóricas
```

```
df$is_attributed <- as.factor(df$is_attributed)
```

```
# Remover colunas não utilizadas
```

```
df$click_time <- NULL
```

```
df$attributed_time <- NULL
```

```

str(df)

## 'data.frame': 100000 obs. of 8 variables:
## $ ip : int 87540 105560 101424 94584 68413 93663 17059 121505 192967 143636 ...
## $ app : int 12 25 12 13 12 3 1 9 2 3 ...
## $ device : int 1 1 1 1 1 1 1 1 2 1 ...
## $ os : int 13 17 19 13 1 17 17 25 22 19 ...
## $ channel : int 497 259 212 477 178 115 135 442 364 135 ...
## $ is_attributed: Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
## $ click_day : int 7 7 7 7 9 9 9 7 8 8 ...
## $ click_hour : int 9 13 18 4 9 1 1 10 9 12 ...

# Analise Exploratoria de Dados

# Carregando os Pacotes
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:data.table':
##
## between, first, last

## The following objects are masked from 'package:stats':
##
## filter, lag

## The following objects are masked from 'package:base':
##
## intersect, setdiff, setequal, union

# Verificar se existem valores ausentes (missing) em cada coluna
# Nenhum valor encontrado
any(is.na(df))

## [1] FALSE

# Analisando dados por agrupamentos
df %>% group_by(ip) %>% tally()

## # A tibble: 34,857 x 2
## ip n
## <int> <int>
## 1 9 1
## 2 10 3
## 3 19 1
## 4 20 4
## 5 25 1
## 6 27 5
## 7 31 1
## 8 33 1
## 9 36 3
## 10 59 3
## # ... with 34,847 more rows

df %>% group_by(ip, click_day) %>% tally()

```

```
## # A tibble: 55,454 x 3
## # Groups:   ip [34,857]
##       ip click_day     n
##   <int>    <int> <int>
## 1     9         7     1
## 2    10         7     2
## 3    10         8     1
## 4    19         8     1
## 5    20         8     3
## 6    20         9     1
## 7    25         7     1
## 8    27         7     1
## 9    27         8     2
## 10   27         9     2
## # ... with 55,444 more rows
```

```
# Adicionando nova coluna no dataframe
```

```
df_count_ip <- df %>%
  count(ip, sort = TRUE, name = "ip_count")
```

```
df <- merge(df, df_count_ip, by=c("ip"))
rm('df_count_ip')
```

```
df_count_ip_day <- df %>%
  count(ip, click_day, sort = TRUE, name = "ip_day_count")
```

```
df <- merge(df, df_count_ip_day, by=c('ip','click_day'))
rm('df_count_ip_day')
```

```
str(df)
```

```
## 'data.frame':   100000 obs. of  10 variables:
## $ ip           : int  10 10 10 1000 100002 100005 100005 100009 100013 100013 ...
## $ click_day    : int  7 7 8 7 8 7 9 8 8 9 ...
## $ app          : int  11 12 18 12 3 2 9 64 3 13 ...
## $ device       : int  1 1 1 1 1 1 1 1 1 1 ...
## $ os           : int  22 19 13 19 41 17 19 18 41 10 ...
## $ channel      : int  319 140 107 178 280 219 232 459 442 477 ...
## $ is_attributed: Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
## $ click_hour   : int  1 7 11 13 2 3 14 9 5 7 ...
## $ ip_count     : int  3 3 3 1 1 2 2 1 2 2 ...
## $ ip_day_count : int  2 2 1 1 1 1 1 1 1 1 ...
```

```
# View(df)
```

```
# Normalizar as variaveis numericas
```

```
cols <- c('ip', 'app', 'device', 'os', 'channel','click_day','click_hour','ip_count', 'ip_day_count')
df[, cols] <- scale(df[, cols])
```

```
# Verificando overfitting dos dados
```

```
# 99.773 registros indicam que o app nao foi baixado
```

```
# 227 registros indicam que o app foi baixado
```

```
table(df$is_attributed)
```

```
##
##      0      1
```

```

## 99773    227
prop.table(table(df$is_attributed))

##
##          0          1
## 0.99773 0.00227

# Feature Selection (Selecao de Variaveis)

# Carregando os Pacotes
library(ROSE)

## Loaded ROSE 0.0-3
library(caret)

## Loading required package: lattice
## Loading required package: ggplot2
## Registered S3 methods overwritten by 'ggplot2':
##   method      from
##   [.quosures   rlang
##   c.quosures   rlang
##   print.quosures rlang

library(e1071)
library(rpart)

# Gerando dados de treino e de teste
splits <- createDataPartition(df$is_attributed, p=0.7, list=FALSE)

# Separando os dados
dados_treino <- df[ splits,]
dados_teste <- df[-splits,]

# Verificando o numero de linhas
nrow(dados_treino)

## [1] 70001
nrow(dados_teste)

## [1] 29999

# Treinando o modelo usando Naive Bayes e fazendo predicoes
## devido ao problema de overfitting, o resultado esta tendencioso
## necessario corrigir o problema de overfitting
modeloNB <- naiveBayes(is_attributed ~. , data=dados_treino)
predNB <- predict(modeloNB, dados_teste)
confusionMatrix(predNB, dados_teste$is_attributed)

## Confusion Matrix and Statistics
##
##          Reference
## Prediction    0    1
##          0 29173   59
##          1   758    9


```

```
##
##          Accuracy : 0.9728
##          95% CI : (0.9709, 0.9746)
##    No Information Rate : 0.9977
##    P-Value [Acc > NIR] : 1
##
##          Kappa : 0.0175
##
##    Mcnemar's Test P-Value : <2e-16
##
##          Sensitivity : 0.97468
##          Specificity : 0.13235
##    Pos Pred Value : 0.99798
##    Neg Pred Value : 0.01173
##          Prevalence : 0.99773
##    Detection Rate : 0.97247
##    Detection Prevalence : 0.97443
##    Balanced Accuracy : 0.55351
##
##    'Positive' Class : 0
##
```

```
# AUC
```

```
roc.curve(dados_teste$is_attributed, predNB)
```

**ROC curve**



```
## Area under the curve (AUC): 0.554
```

```
# Resolvendo problema de Overfitting usando pacote ROSE
```

```
#over sampling
```

```
dados_treino_new <- ROSE(is_attributed ~ . , data=dados_treino)$data
```


```
table(dados_treino_new$is_attributed)
```

```
##
##      0      1
## 35177 34824
prop.table(table(dados_treino_new$is_attributed))

##
##      0      1
## 0.5025214 0.4974786
# Treinando um novo modelo com os novos dados de treino
modeloNB_v2 <- naiveBayes(is_attributed ~ . , data=dados_treino_new)
predNB_v2 <- predict(modeloNB_v2, dados_teste)
confusionMatrix(predNB_v2, dados_teste$is_attributed)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction      0      1
##      0 27456      37
##      1  2475      31
##
##              Accuracy : 0.9163
##              95% CI : (0.9131, 0.9194)
##      No Information Rate : 0.9977
##      P-Value [Acc > NIR] : 1
##
##              Kappa : 0.0198
##
##      Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.91731
##              Specificity : 0.45588
##              Pos Pred Value : 0.99865
##              Neg Pred Value : 0.01237
##              Prevalence : 0.99773
##              Detection Rate : 0.91523
##              Detection Prevalence : 0.91646
##              Balanced Accuracy : 0.68660
##
##      'Positive' Class : 0
##
# AUC
roc.curve(dados_teste$is_attributed, predNB_v2)
```

## ROC curve



```
## Area under the curve (AUC): 0.687
```

```
#AUC ROSE
```

```
ROSE.holdout <- ROSE.eval(is_attributed ~ .,
                          data = dados_treino_new,
                          learner = rpart,
                          method.assess = "holdout",
                          extr.pred = function(obj) obj[,2])

ROSE.holdout
```

```
##
```

```
## Call:
```

```
## ROSE.eval(formula = is_attributed ~ ., data = dados_treino_new,
##          learner = rpart, extr.pred = function(obj) obj[, 2], method.assess = "holdout")
##
```

```
## Holdout estimate of auc: 0.846
```

```
# Análise de Correlação
```

```
# Carregando os Pacotes
```

```
library(corrplot)
```

```
## corrplot 0.84 loaded
```

```
library(corrgram)
```

```
## Registered S3 method overwritten by 'seriation':
```

```
##   method      from
```

```
## reorder.hclust gclus
```

```
##
```

```
## Attaching package: 'corrgram'
```

```
## The following object is masked from 'package:lattice':
##
##   panel.fill
```

```
# obtendo somente as colunas numericas
colunas_numericas <- sapply(dados_treino_new, is.numeric)
colunas_numericas
```

```
##           ip      click_day      app      device      os
##      TRUE      TRUE      TRUE      TRUE      TRUE
## channel is_attributed click_hour ip_count ip_day_count
##      TRUE      FALSE      TRUE      TRUE      TRUE
```

```
# Filtrando as colunas numericas para correlacao
data_cor <- cor(dados_treino_new[,colunas_numericas])
data_cor
```

```
##           ip      click_day      app      device
## ip      1.000000000  0.188812587  0.105519985  0.004707135
## click_day 0.188812587  1.000000000 -0.046530254 -0.031124922
## app      0.105519985 -0.046530254  1.000000000  0.053060671
## device   0.004707135 -0.031124922  0.053060671  1.000000000
## os       0.119795925  0.028070439  0.132279079  0.329260697
## channel  -0.116477683  0.004050537 -0.093331558 -0.015940436
## click_hour 0.026053281 -0.208323592  0.026232579  0.004416267
## ip_count  -0.176095858 -0.081155777 -0.003287766  0.002044292
## ip_day_count -0.161268845 -0.052111784 -0.030515931  0.004682263
##           os      channel      click_hour      ip_count
## ip      0.119795925 -0.116477683  0.026053281 -0.176095858
## click_day 0.028070439  0.004050537 -0.208323592 -0.081155777
## app      0.132279079 -0.093331558  0.026232579 -0.003287766
## device   0.329260697 -0.015940436  0.004416267  0.002044292
## os       1.000000000 -0.017273315  0.014947085 -0.012736725
## channel  -0.017273315  1.000000000  0.056284018 -0.013373677
## click_hour 0.014947085  0.056284018  1.000000000  0.145328667
## ip_count  -0.012736725 -0.013373672  0.145328667  1.000000000
## ip_day_count -0.008899855 -0.000293483  0.123115635  0.704119166
##           ip_day_count
## ip      -0.161268846
## click_day -0.052111783
## app      -0.030515931
## device   0.004682267
## os       -0.008899855
## channel  -0.000293483
## click_hour 0.123115635
## ip_count  0.704119166
## ip_day_count 1.000000000
```

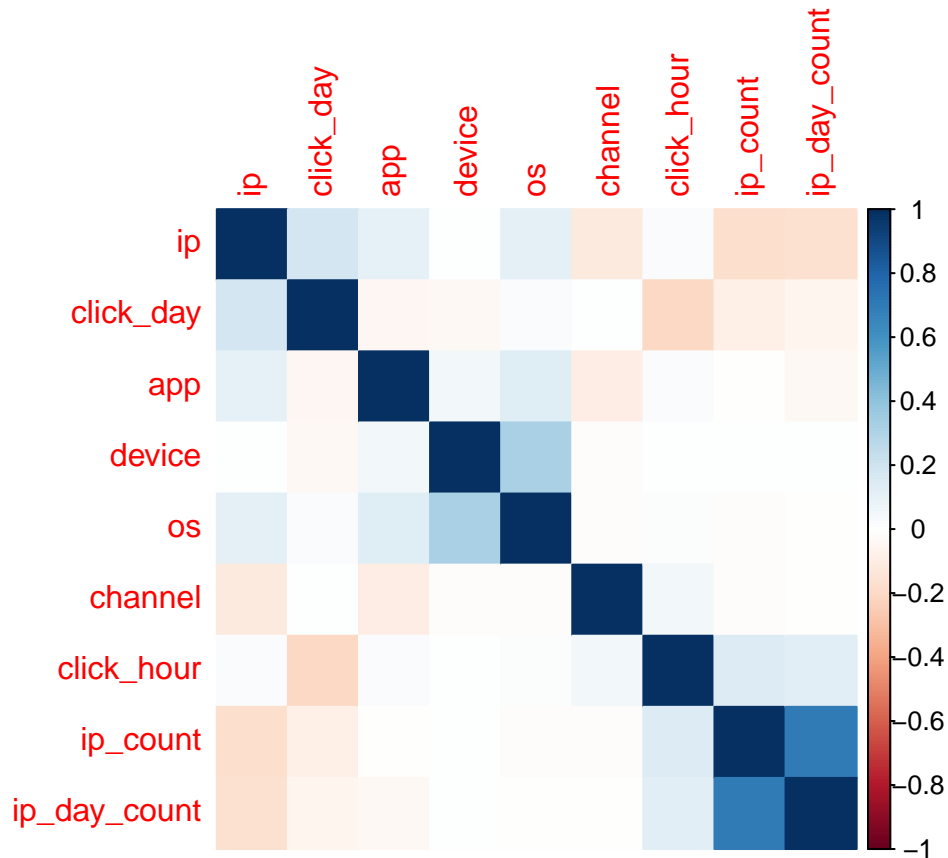
```
head(data_cor)
```

```
##           ip      click_day      app      device      os
## ip      1.000000000  0.188812587  0.10551999  0.004707135  0.11979592
## click_day 0.188812587  1.000000000 -0.04653025 -0.031124922  0.02807044
## app      0.105519985 -0.046530254  1.00000000  0.053060671  0.13227908
## device   0.004707135 -0.031124922  0.05306067  1.000000000  0.32926070
## os       0.119795925  0.028070439  0.13227908  0.329260697  1.00000000
```




```
## channel -0.116477683 0.004050537 -0.09333156 -0.015940436 -0.01727331
##          channel click_hour ip_count ip_day_count
## ip -0.116477683 0.026053281 -0.176095858 -0.1612688446
## click_day 0.004050537 -0.208323592 -0.081155777 -0.0521117839
## app -0.093331558 0.026232579 -0.003287766 -0.0305159312
## device -0.015940436 0.004416267 0.002044292 0.0046822627
## os -0.017273315 0.014947085 -0.012736725 -0.0088998554
## channel 1.000000000 0.056284018 -0.013373677 -0.0002934833
```

```
# Criando um corrplot
corrplot(data_cor, method = 'color')
```



```
# Criando um corrgram
corrgram(dados_treino_new, order=TRUE, lower.panel = panel.shade,
         upper.panel = panel.pie, text.panel = panel.txt)
```



```
# Cria um modelo preditivo usando randomForest
```

```
# Carregando os Pacotes
```

```
library(randomForest)
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
## margin
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
## combine
```

```
# Cria o modelo preditivo usando randomForest
```

```
modeloRF <- randomForest(is_attributed ~ .,
                          data = dados_treino_new,
                          ntree = 40,
                          nodesize = 5)
```

```
print(modeloRF)
```

```
##
```

```
## Call:
```

```
## randomForest(formula = is_attributed ~ ., data = dados_treino_new, ntree = 40, nodesize = 5)
```

```
## Type of random forest: classification
```

```
## Number of trees: 40
```

```
## No. of variables tried at each split: 3
```

```
##
```

```
## OOB estimate of error rate: 9.45%
```

```
## Confusion matrix:
```

```
##          0          1 class.error
## 0 32276 2901 0.08246866
## 1 3715 31109 0.10667930

# Previsões com um modelo de classificação baseado em randomForest

# Gerando previsões nos dados de teste
previsoes <- data.frame(observado = dados_teste$is_attributed,
                        previsto = predict(modeloRF, newdata = dados_teste))

# Visualizando o resultado
# View(previsoes)

# Calculando a Confusion Matrix em R

# Carregando os Pacotes
library(ROCR)


## Loading required package: gplots

##
## Attaching package: 'gplots'

## The following object is masked from 'package:stats':
##
##      lowess

# Gerando as classes de dados
class1 <- predict(modeloRF, newdata = dados_teste, type = 'prob')
class2 <- dados_teste$is_attributed

# Gerando a curva ROC
pred <- prediction(class1[,2], class2)
perf <- performance(pred, "tpr", "fpr")
plot(perf, col = rainbow(10))
```



```
# Gerando Confusion Matrix com o Caret
# Dataframes com valores observados e previstos
previsoes_v2 <- data.frame(observado = dados_teste$is_attributed,
                           previsto = predict(object = modeloRF, newdata = dados_teste))

confusionMatrix(previsoes_v2$observado, previsoes_v2$previsto)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0 29577   354
##           1   41    27
##
##           Accuracy : 0.9868
##           95% CI : (0.9855, 0.9881)
##           No Information Rate : 0.9873
##           P-Value [Acc > NIR] : 0.7738
##
##           Kappa : 0.1169
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.99862
##           Specificity : 0.07087
##           Pos Pred Value : 0.98817
##           Neg Pred Value : 0.39706
##           Prevalence : 0.98730
##           Detection Rate : 0.98593
##           Detection Prevalence : 0.99773
##           Balanced Accuracy : 0.53474
##
##           'Positive' Class : 0
```

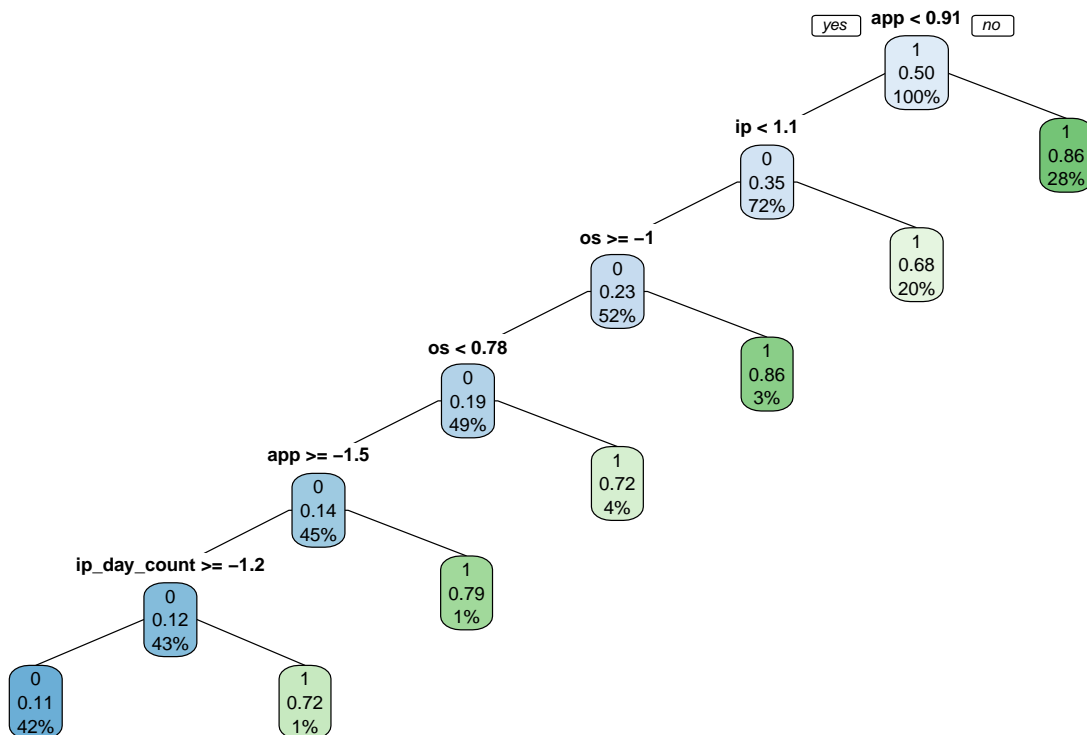
```
##
# Otimizando o Modelo preditivo

# Carregando os Pacotes
library(rpart.plot)

# Criando uma Cost Function
Cost_func <- matrix(c(0, 1.5, 1, 0), nrow = 2, dimnames = list(c("1", "2"), c("1", "2")))

# Criando o Modelo usando rpart
modeloTree <- rpart(is_attributed ~ .,
                     data = dados_treino_new,
                     method = 'class',
                     parms = list(loss = Cost_func))

# Plot do modelo
rpart.plot(modeloTree, fallen.leaves = FALSE, type = 1)
```



```
# Analisando Confusion Matrix
pred.tree <- predict(modeloTree, type = "class")
confusionMatrix(pred.tree, dados_treino_new$is_attributed)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 26467 3209
##           1  8710 31615
##
##           Accuracy : 0.8297
```

```

##              95% CI : (0.8269, 0.8325)
##      No Information Rate : 0.5025
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.6597
##
##      McNemar's Test P-Value : < 2.2e-16
##
##              Sensitivity : 0.7524
##              Specificity : 0.9079
##              Pos Pred Value : 0.8919
##              Neg Pred Value : 0.7840
##              Prevalence : 0.5025
##              Detection Rate : 0.3781
##      Detection Prevalence : 0.4239
##              Balanced Accuracy : 0.8301
##
##      'Positive' Class : 0
##
# Modelo usando Naive Bayes com os novos dados de treino
modeloNB_v2 <- naiveBayes(is_attributed ~ . , data=dados_treino_new)
predNB_v2 <- predict(modeloNB_v2, dados_teste)
confusionMatrix(predNB_v2, dados_teste$is_attributed)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction      0      1
##              0 27456    37
##              1  2475    31
##
##              Accuracy : 0.9163
##              95% CI : (0.9131, 0.9194)
##      No Information Rate : 0.9977
##      P-Value [Acc > NIR] : 1
##
##              Kappa : 0.0198
##
##      McNemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.91731
##              Specificity : 0.45588
##              Pos Pred Value : 0.99865
##              Neg Pred Value : 0.01237
##              Prevalence : 0.99773
##              Detection Rate : 0.91523
##      Detection Prevalence : 0.91646
##              Balanced Accuracy : 0.68660
##
##      'Positive' Class : 0
##
# Modelo usando RandomForest com os novos dados de treino
modeloRF_v2 <- randomForest(is_attributed ~ . ,

```

```

        data = dados_treino_new,
        ntree = 100,
        nodesize = 5)
predRF_v2 <- predict(modeloRF_v2, dados_teste)
confusionMatrix(predRF_v2, dados_teste$is_attributed)

```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction      0      1
```

```
##           0 29565    41
```

```
##           1   366    27
```

```
##
```

```
##           Accuracy : 0.9864
```

```
##           95% CI : (0.9851, 0.9877)
```

```
## No Information Rate : 0.9977
```

```
## P-Value [Acc > NIR] : 1
```

```
##
```

```
##           Kappa : 0.1137
```

```
##
```

```
## McNemar's Test P-Value : <2e-16
```

```
##
```

```
##           Sensitivity : 0.9878
```

```
##           Specificity : 0.3971
```

```
## Pos Pred Value : 0.9986
```

```
## Neg Pred Value : 0.0687
```

```
## Prevalence : 0.9977
```

```
## Detection Rate : 0.9855
```

```
## Detection Prevalence : 0.9869
```

```
## Balanced Accuracy : 0.6924
```

```
##
```

```
## 'Positive' Class : 0
```

```
##
```